



# Echo State Networks for Time Series Data

Open Data Science Conference, Boston

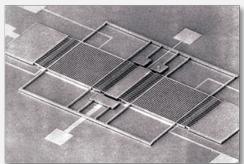
Teal Guidici, PhD

April 16<sup>th</sup>, 2020

# About Draper



Apollo Guidance Computer



Microelectromechanical System (MEMS)



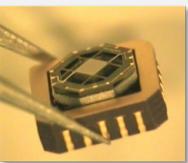
First Algebraic Compiler



Precision Guided Munitions



Tissue Engineering



Precision Timing



Ballistic Missile Guidance



Digital Fly-by-Wire



Shuttle Docking



SPIRE Jr.



Mark 14 Gunsight



Precision Airdrop



Engineering the World Around You

Independent not-for-profit research and development company

- Headquartered in Cambridge, MA
- 1500+ employees, \$520 million revenue
- Serving the nation since 1932
- Independent organization since 1973

# Outline

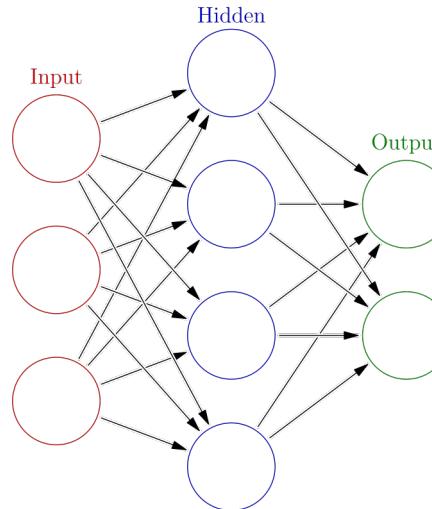
---

- Neural Networks in general
- Echo State Networks in detail
- Demo: [https://github.com/tguidici/ODSC\\_2020\\_ESN](https://github.com/tguidici/ODSC_2020_ESN)
- Example: ESN for finance

# Neural Networks (NN)

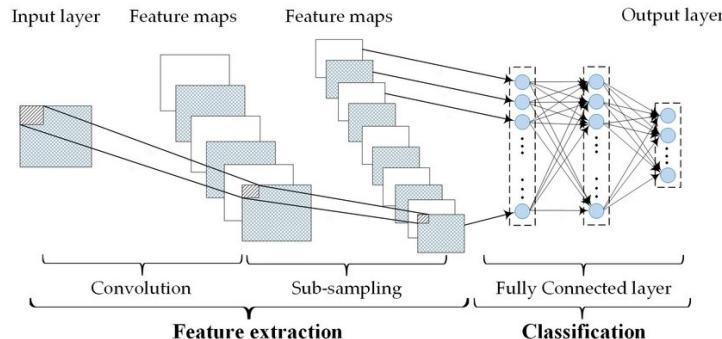
---

- Loosely (!) modeled on human/animal brains
  - Many simple processing nodes
  - Densely connected (usually)
  - “Learn” to perform a task by analyzing training examples
- First proposed in 1944; on a boom-bust cycle of interest
  - Dropped out of favor ~ 1969
  - Resurgence in 1980’s!
  - Disinterest in 2000’s!
  - ALL THE INTEREST IN 2010’s!!



- Circular nodes represent artificial neurons
- Arrows illustrate connections between output of one neuron to input of another

# Convolutional vs Recurrent Neural Nets

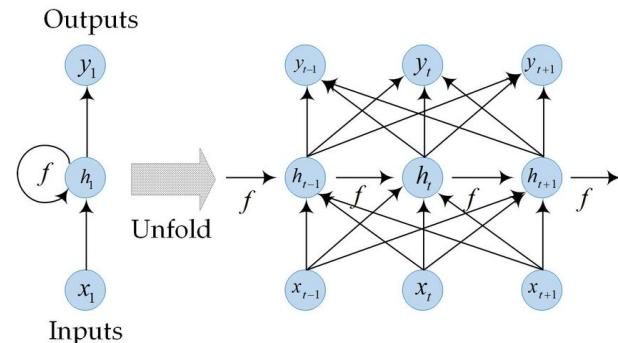


**Convolutional Neural Net**



Spatial structure matters  
eg: image classification

Am I a CAT or a COWPOKE?

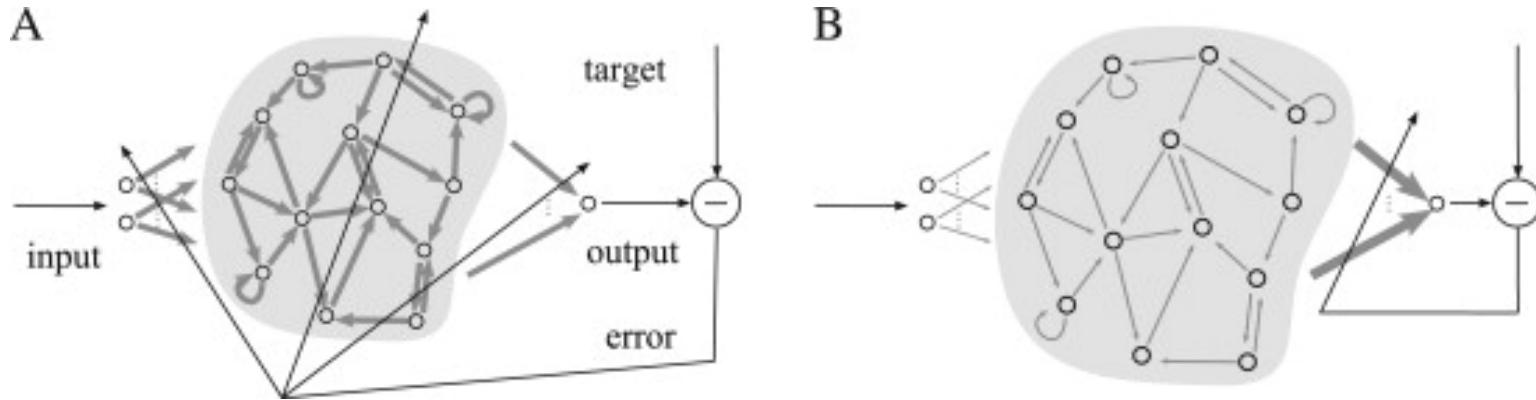


**Recurrent Neural Net (RNN)**

Sequential structure matters  
eg: weather prediction

It rained the last two days, will it rain today?

# Reservoir Computing – a better RNN



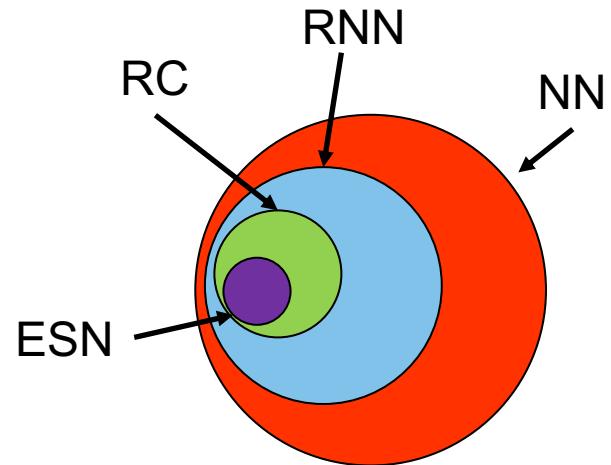
- A. Traditional RNN: all connection weights adapted
  - Computationally intensive, hard to tune
- B. Reservoir Computing (RC): only output weights are adapted
  - Fast! Accurate! Still challenging to tune

Bold arrows  
indicate trained  
weights

# Echo State Networks (ESN) – introduction!

---

- ESNs are a recurrent neural network
  - Captures **dynamic behavior** of a time series or sequential data.
  - Internal state (**memory**) to process sequences of inputs and remember context
- ESNs are also type of reservoir computing
  - **Faster training** by only learning the output weights vs. back propagation approaches
  - Captures more complex interactions by manipulating topologies and neurons within network
  - Non-linear transformation of inputs into higher dimension of reservoir
  - Output weights are trained ~ “what linear combination of reservoir elements best approximates my target?”



ESNs are a type of Reservoir Computing, which is a subset of RNNs, which are a type of Neural Network

# ESN: Overview

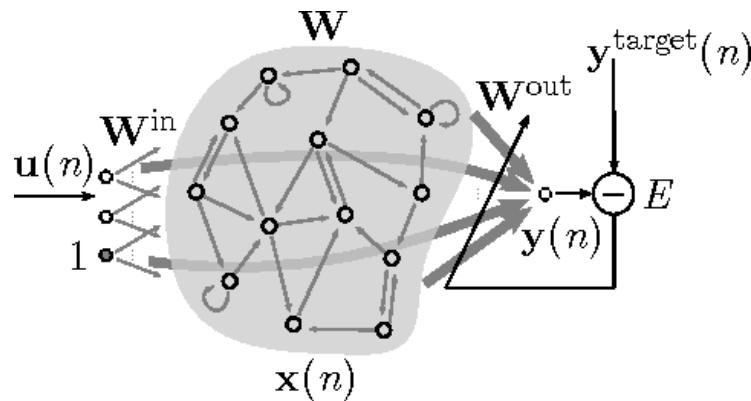


Fig. 1: An echo state network.

For each time step n

- $u(n)$ : input
- $x(n)$ : reservoir activations
- $y(n)$ : network output
- $y^{\text{target}}(n)$ : desired output

**Randomly Initialized & fixed**

- $W^{\text{in}}$ : input weights
  - Linearly map input into reservoir
- $W$ : recurrent weights
  - How are the reservoir neurons connected to each other?

**Learned parameters**

- $W^{\text{out}}$ : output weights
  - Linearly map  $(u(n), x(n))$  to  $y(n)$
  - Minimize Error( $y^{\text{target}}(n)$ ,  $y(n)$ )

# ESN: Overview

---

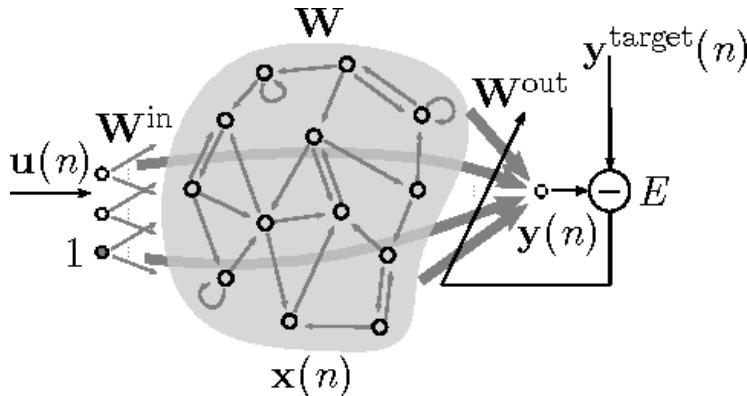


Fig. 1: An echo state network.

## How does it work?

1. **Setup the network!** : Generate a large random reservoir RNN with
  - Input weights  $W^{\text{in}}$
  - Recurrent weights  $W$
  - Leak rate  $\alpha$
2. **Map to higher dimension!**: Run input  $u(n)$  through network and collect activation states  $x(n)$ 
  - Non-linear high dimensional transform!
3. **Optimize!**: Compute  $W^{\text{out}}$  such that  $W^{\text{out}}$  minimizes Error( $y(n)$ ,  $y^{\text{target}}(n)$ )
  - The only training we need!
  - Error minimized across all time points
4. **Predict something new!**: Trained network can be used to compute predicted  $y(n)$  from new  $u(n)$

# Definitions and update equations

---

$W^{in} \in \mathbb{R}^{N_x \times (1+N_u)}$ : input weights

$W \in \mathbb{R}^{N_x \times N_x}$  : recurrent weights

$\alpha \in (0,1]$ : leak rate

$x(n) \in \mathbb{R}^{N_x}$  : neuron activations

$\tilde{x}(n) \in \mathbb{R}^{N_x}$  : neuron updates

$$x(n) = (1 - \alpha)x(n - 1) + \alpha\tilde{x}(n)$$

$$\tilde{x}(n) = \tanh(W^{in}[1; u(n)] + Wx(n - 1))$$

# Key Reservoir Parameters: Input Weights

---

$$W^{in} \in \mathbb{R}^{N_x \times (1+N_u)}$$

$$W_{ij}^{in} \sim Unif(-a, a)$$

Tune/scale input weights by varying  $a$

- Generally drawn from a Uniform distribution, but not required!
- Input scaling determines how nonlinear reservoir responses are
  - $a < 1$  (generally)
- Scaling input weights helps remove the influence of outliers.
- Different sets of inputs can be scaled differently if the corresponding portions of  $u(n)$  contribute differently to task

# Key Reservoir Parameters: Spectral Radius

---

$$W \in \mathbb{R}^{N_x \times N_x}$$

$$\rho(W) = \text{spectral radius of } W$$

- Spectral radius = eigenvalue of  $W$  with greatest magnitude
- $W$  defines which reservoir neurons are connected and how
- $\rho(W) < 1$  ensures echo state property. But not always necessary.
- $y(n)$  requires extensive historical knowledge? Large  $\rho(W)$
- $y(n)$  depends on more recent history? Small  $\rho(W)$
- Practical approach: Select  $\rho(W)$  that maximizes accuracy, using  $\rho(W)=1$  as initial starting value

# Key Reservoir Parameters: Leak Rate

---

$\alpha \in (0, 1]$ : leak rate

$x(n) = (1 - \alpha)x(n - 1) + \alpha\tilde{x}(n)$ : neuron activations

$\tilde{x}(n) = \tanh(W^{in}[1; u(n)] + Wx(n - 1))$ : neuron updates

- Captures temporal dynamics of system
- "speed of reservoir update dynamics discretized in time"
- Goal: set  $\alpha$  to match the speed of the dynamics of  $u(n)$  and/or  $y^{\text{target}}(n)$
- Different components of  $u(n)$  have different speed dynamics? No problem!

$$\alpha \in \mathbb{R}^{N_x}$$

# Other reservoir parameters

---

**Can generally be set to default/“reasonable” values**

- Sparsity of the reservoir (connectivity)
  - Reservoirs should be very sparsely connected, regardless of size.
    - Potentially each neuron only connected to 10 others!
    - This sparsity is what makes it fast! fast! fast!
- Size of the reservoir (number of neurons)
  - Pick as many as you think you'll need computationally
  - But you probably need fewer than you think (10's, not 100's)
- Distribution used to generate weights for  $W$ 
  - Uniform distribution is probably good enough.
  - Can also be used for generating weights for  $W^{\text{in}}$

# More choices – Which error?

---

Output weights,  $W^{out}$ , computed to minimize error....but which error?

$$W^{out} = \underset{W^{out}}{\operatorname{argmin}} E(y, y^{target})$$

- Simplest: Mean Squared Error
- Potential for overfitting or feedback instability? Ridge Regression
- No overfitting and lots of compute power? Pseudoinverse
- Some time points more (less) important than others? Regression Weighting
  - Weight more (less) heavily error for important timepoints

D R A P E R®

# ESN DEMO

# DEMO

---

DEMO can be found at: [https://github.com/tguidici/ODSC\\_2020\\_ESN](https://github.com/tguidici/ODSC_2020_ESN)

D R A P E R®

# ESN for Finance

# Echo State Networks for Market Modeling

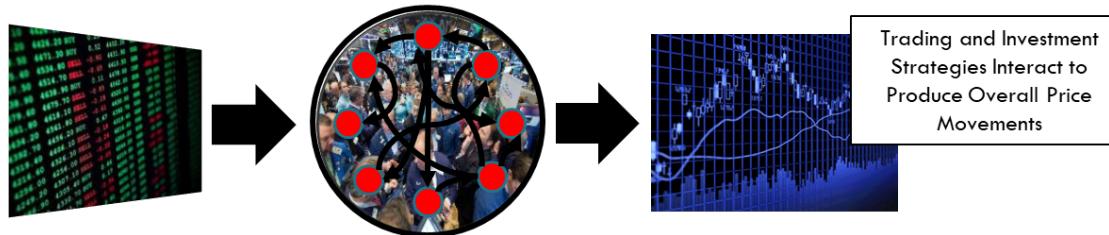
1. The market is a complex system
  - Large system of non linearly interacting information and behaviors.
  - Temporal/sequential information crucial (“memory”)
2. Echo State Networks are uniquely capable of modeling non-linear, interacting dynamics with temporal dependencies.

**Theory** : The Stock Market is a Network of Interacting Strategies

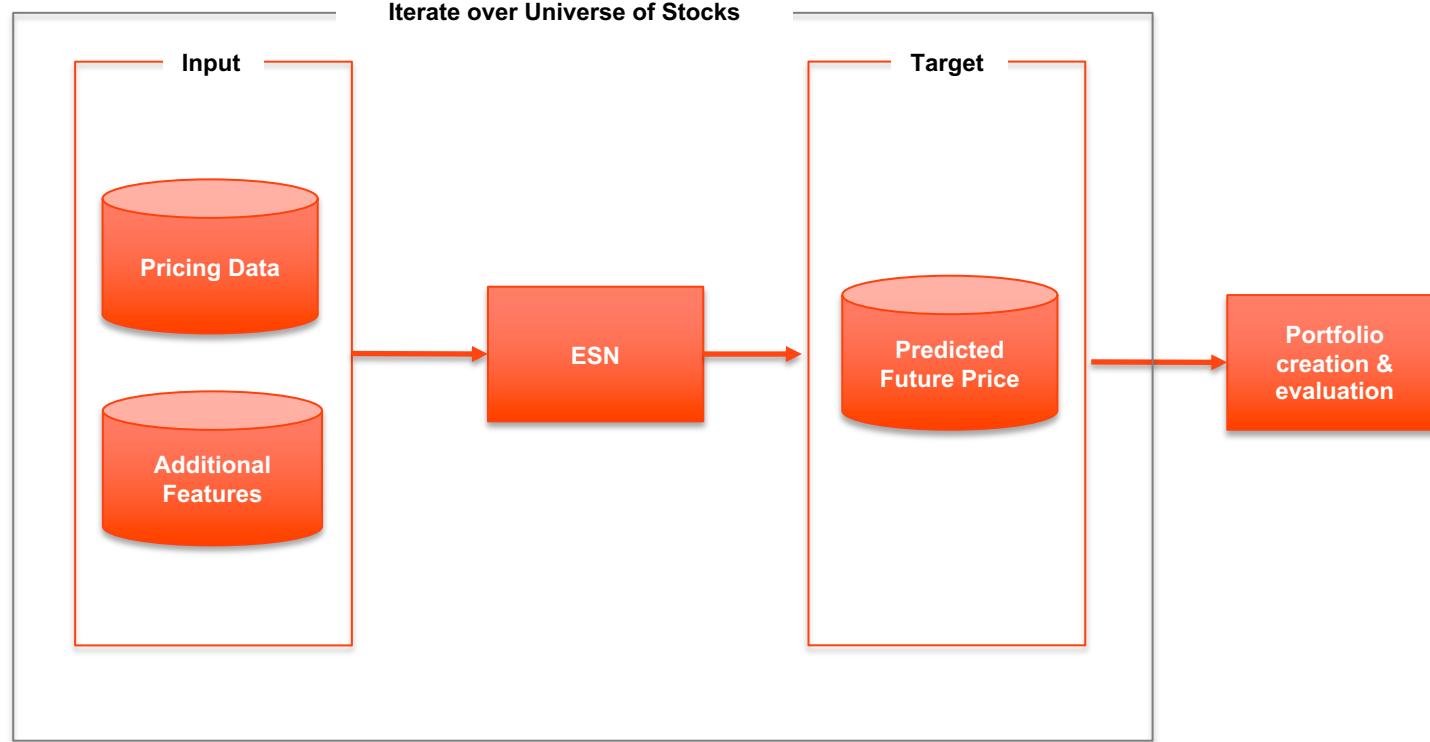
**Input** – Market, company news, external news, world events, opinions

**Computation** – Complex system dynamics of strategies

**Output** – Predicted stock price movement and corresponding volume



# ESN for Market Modeling pipeline



# Test Setup: parameter tuning & data leakage

---

**Data Leakage:** when your model knows more than it should

	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
<b>Development Test</b>										
Train										
Validate										
Out of Sample Test										
<b>Final Test #1</b>										
Train										
Validate										
Out of Sample Test										

- All development tests in the pre-2010 timeframe
  - True out-of-sample tests in the post-2010 timeframe
  - Final tests were executed a single time using parameters determined during development
-

# ESN for Market Modeling: data leakage

---

## Data Leakage: when your model knows more than it should

What we did well:

- Held out “final” datasets for test of the entire pipeline
- Separate training and validation datasets for hyperparameter tuning
- Temporal buffer to separate validation and test datasets

What we didn’t do well:

- Optional PCA featurization used data before splitting into train/validate/test
  - features in training data “knew” something about variance of test data  
(Caught this error before final tests!)

# Automated hyperparameter tuning

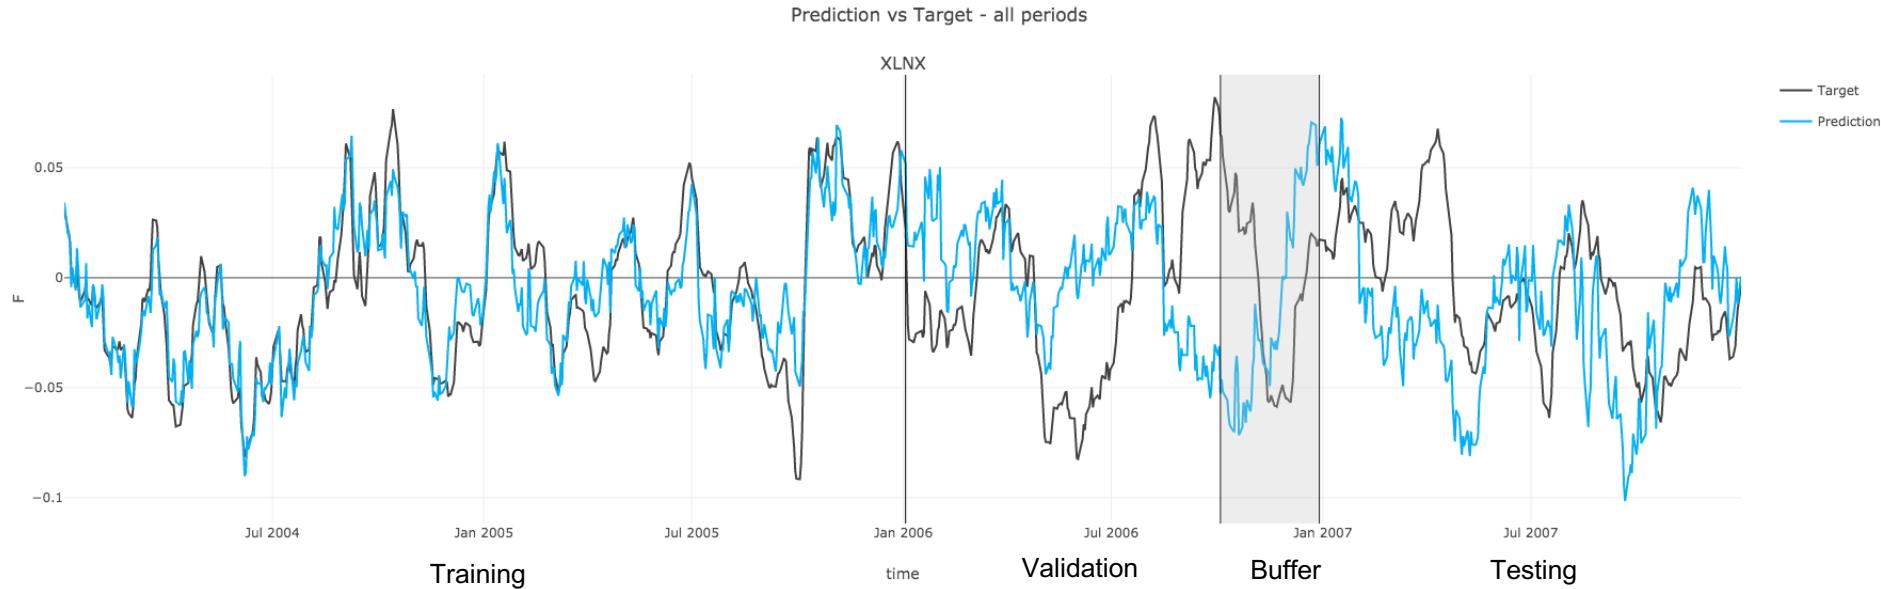
---

**ESNs have many hyperparameters – how to pick good values?**

We automated the pipeline to iterate over a "reasonable" range for select parameters, pick values with lowest error, on average.

- Training:
  - Spectral radius: tune over reservoir
  - Input scaling: tune over inputs
- Validation:
  - Leak rate: tune over temporal dynamics of input

# ESN for Finance – Results



Portfolios constructed based on ESN predictions were robust and performed well

# Where to next?

---

- Ensemble or on-line training
  - Data generating process not stationary? Model needs to adapt
- Output feedback
  - Trained readouts fed back into reservoir to change dynamics
- Semi-structured network topologies
  - Groups of neurons linked to specific inputs
- Different error metrics
  - Some types of error matter more than others (eg: when predicting counts, output should not be negative....)